

Project 2 – Distributed Shell

Project 2 – Distributed Shell

Due: Thursday, April 5th

Motivation

- Shell is core operating systems concept
- Executing commands on remote machine is core distributed systems concept
 - Computing in the “cloud” (except in this case, server address is known, not a service) with results returned to user

→ *Distributed Shell*

Distributed Shell

■ Client-Server

- Server on “known” port for client

■ Non-interactive

- Command line args
- `get-opt.c`
- Authentication built-in

■ Uses TCP sockets


- `listen.c` and `talk.c`

■ Security

- Authentication (next slide)

■ Can handle multiple requests

- `fork()`
- Concurrent server



Simple Authentication

- **Client connects to server, sending in user-name (not password)**
 - Note, typically name used to lookup unique data (salt) stored with each name that further randomizes hash salt – here, just a check that valid
- **Server responds by sending back unique random number**
 - Use `rand()`
 - Changes each invocation (provide different random seed each run, using `srand()`)
- **Client encrypts using password and number as key**
 - Here, password can be “hard coded” into client
 - Use `crypt()`
- ...

Simple Authentication

- ...
- Client sends hashed value to server
- Server encrypts with user's password and same number as key
 - Server “knows” password (e.g., externally exchanged)
 - Here, password can be “hard coded” into server
- Server compares two hashed/encrypted values, if same then ok
 - Otherwise, return error message and exit
- Note, can do authentication either before or after `fork()`
 - Question: pluses and minuses for each?

Sample

Server

```
YourVM% ./server  
./server activating.  
port: 4513  
dir: /home/claypool/dsh Socket  
created!  
Accepting connections.  
  
Connection request received.  
forked child  
received: john  
password ok  
command: ls  
executing command...
```

Client

```
OtherVM% ./dsh -c "ls" -s cccwork3  
Makefile  
client.c  
dsh  
dsh.c  
index.html  
server  
server.c  
server.h  
sock.c  
sock.h
```

Hints

- **Build shell independent of server**
- **Get basic connection working without shell**
 - Socket help: `listen-tcp.c` and `talk-tcp.c`
 - Socket slides (next deck)
- **Shell help**
 - `fork.c`
 - `execl.c`
 - `fork()`, `execve()`, `getopt()`, `strtok()`, `dup2()`
- **Use a Makefile**
 - Web page's is simple
- **Use man pages for functionality, return codes**
- **Beware of zombies!**



Debug locally

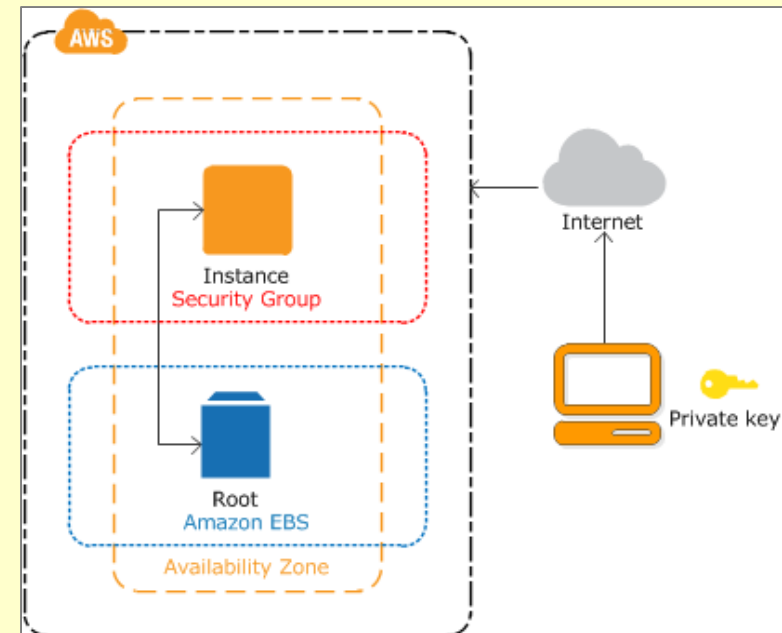
- Do all of this on a pair of course virtual machines
- Show that you can write a command locally that will execute remotely
 - On your teammate's VM
- When that works ...

Cloud Computing – Amazon




Setup Amazon Elastic Compute Cloud (EC2)

1. **Sign up (free, 1 year)**
2. **Launch Linux VM**
 - Copy security keys
3. **Connect (e.g., putty)**
4. **Clean up (when done)**




Cloud Computing – Amazon

- Will have “cloud” Linux host to admin
 - Note IP of instance when starting
 - Install software
 - `gcc` (or `g++`)
 - `make`
 - `sysbench` (for performance)
 - Typically: `sudo apt-get install {name}`
 - Copy server source (e.g., `server.c` and `Makefile`) to cloud host
 - Rebuild
 - Re-test
 - Debug (if necessary)
- (See Project Description for links to documentation)



Experiments – Network Performance

- **Latency of connection (milliseconds)**
 - Includes authentication
 - Force call to server so that server does not do `exec ()`
- **Maximum throughput (b/s)**
 - Transfer large file
 - `redirect > to file` else `stdout` may be bottleneck
- **Multiple runs**



Experiments –

CPU and File I/O Performance

■ Sysbench for CPU

```
sysbench --test=cpu --cpu-max-prime=20000 run
```

- *total time: 23.8724s*

■ Sysbench for File I/O

```
sysbench --test=fileio --file-total-size=16G prepare
```

```
sysbench --test=fileio --file-total-size=16G --file-test-mode=rndrw --init-rng=on --max-time=30 --max-requests=0 run
```

- *Read 9.375Mb ...(53.316Kb/sec)*

■ Clean up!


```
sysbench --test=fileio --file-total-size=16G cleanup
```

■ Local (own virtual machine) vs remote (cloud)

Model

$$\text{Local_CPU} + \text{Local_File_I/O} = n * \text{Network} + \text{Remote_CPU} + \text{Remote_File_I/O}$$

- Simple, algebraic model
- Compute when (for data transfer) cloud computing gives better performance (reduced time for task)
- Solve for n



Writeup

■ ***Design*** - describe your experiments

- Programs/scripts (pseudo-code)
- Number of runs
- Data recording method
- System conditions
 - `lscpu` CPU information (also `/proc/cpuinfo`)
 - `lsblk` block device information (also `df/mount | column -t`)
 - `free -h` available memory
 - `uname -a` OS version (also `/proc/version`)
- Other ...

■ ***Results*** - depict your results *clearly*

- Graph (see previous slide)
- Table (at least *mean* and *standard deviation*)

■ ***Analysis*** - interpret results

- Describe what results mean

Hand In

- **Source code package**
 - `server.c`, `client.c`
 - Support files (`*.h`)
 - `Makefile`
 - `README.txt` detailing building and using
- **Experimental writeup**
- **Clean and zip**
- **Submit via Instruct Assist**

Grading

- Basic shell program (15 points)
- Basic Client-Server communication program (15 points)
- Proper re-direction of standard output (10 points)
- Proper error checking of system/socket calls (10 points)
- Proper authentication (10 points)
- Proper closing of sockets in relation to fork (5 points)
- Proper handling of zombies (5 points)
- Amazon EC2 setup (15 points)
- Experiment Design (6 points)
- Experiment Results (5 points)
- Experiment Analysis (4 points)

See [rubric](#) on Web page, too